We will apply Logistic Regression on Pima In

```python
import numpy as np #linear algebra library of Python
import pandas as pd # build on top of numpy for data analysis, data manipulation and d
import matplotlib.pyplot as plt #plotting library of Python
```

Now let's mount Google drive so that we can upload the diabetes.csv file. You can find the code in the 'Code s

```python
from google.colab import drive
drive.mount('/content/gdrive')
```

⤷

First thing that we do is take a look at the shape of the dataframe (df.shape) and take a look at first 5 lines th

```python
df=pd.read_csv('/content/gdrive/My Drive/Colab Notebooks/diabetes.csv') #import file f
df.head() #shows first 5 lines including column namesdf.shape # number of rows and col
```

⤷

```python
df.shape # provides # rows and # columns of the dataframe df - 768 rows and 9 columns
```

⤷

Now we will assess if the dataset has the same proportion of diabetes vs. non-diabetes cases. At the same t dataset we note that woman #2 has a skin thickness of zero and this is not realistic. It leads us to believe tha was available. This does not apply to columns columns 1 and 9 for obvious reasons.

We use a trick to count the non-zero values of the columns. We convert the data type of the dataframe df to t values to false=0 and all other entries to true=1 . We subsequently add up all True entries per column.

```python
df.astype(bool).sum(axis=0) # counts the number of non-zeros for each column while act
```

⤷

The dataframe is unbalanced as we have 268 ones (diabetes) and thus 500 zeros (no diabetes).

The easiest option could be to eliminate all those patients with zero values, but in this way we would eliminat

Another option is to calculate the median value for a specific column and substitute the zero values for the co

```python
median_BMI=df['BMI'].median()
df['BMI']=df['BMI'].replace(to_replace=0, value=median_BMI)

median_BloodPressure=df['BloodPressure'].median()
df['BloodPressure']=df['BloodPressure'].replace(to_replace=0, value=median_BloodPressu

median_Glucose=df['Glucose'].median()
df['Glucose']=df['Glucose'].replace(to_replace=0, value=median_Glucose)

median_SkinThickness=df['SkinThickness'].median()
df['SkinThickness']=df['SkinThickness'].replace(to_replace=0, value=median_SkinThickne

median_Insulin=df['Insulin'].median()
df['Insulin']=df['Insulin'].replace(to_replace=0, value=median_Insulin)
```

```python
df.head() #shows first 5 lines including column names
```

⊐→

The skin thickness of woman #2 is now 23 (median of that column)

Let's create numpy arrays, one for the features (X) and one for the label (y)

```python
X=df.drop('Outcome', 1).values #drop 'Outcome' column but you keep the index column
y=df['Outcome'].values
```

We import the train_test_split function from sklearn to split the arrays or matrices into random train and test s

Parameters:

test_size : in our case 20% (default=0.25)

random_state: is basically used for reproducing your problem the same every time it is run. If you do not use a
the split you might get a different set of train and test data points and will not help you in debugging in case y
number does not matter

stratify : array-like or None (default=None) If the number of values belonging to each class are unbalanced, us
basically asking the model to take the training and test set such that the class proportion is same as of the w

```python
from sklearn.model_selection import train_test_split #method to split training and tes
X_train, X_test, y_train, y_test=train_test_split(X, y, test_size=0.2, random_state=42
```

```python
print(X_train)
```

```python
print(X_test)
```

Now we are ready to use Logistic Regression

```python
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)
score = model.score(X_test, y_test)
print("Accuracy:", score)
```

Now we will look at other classification KPIs that we discussed in our lessons: Confusion Matrix, ROC, AUC, F

```
from sklearn.metrics import confusion_matrix
y_pred=model.predict(X_test)
confusion_matrix(y_test,y_pred)
```

⮕

Classifier not so good: true positives=25, true negatives=86, false positives=14 and false negatives=29. We c
this case as we would tell a women that she is not diabetic whereas she actually is diabetic. One option to re
but this will increase the amount of False Positives (FPs) as we have seen in lesson 3. Recall in our case is Th
TP/(TP+FP)=64%

ROC (Receiver Operating Characteristic) curve

It is a plot of Recall vs. False Positive Rate (FPR) for the different possible thresholds of the classifier. It show
closer the curve follows the left-hand border and then the top border of the ROC space, the more accurate the
diagonal of the ROC space, the less accurate the test. The area under the curve is a measure of test accuracy

```
from sklearn.metrics import roc_curve
y_pred_proba=model.predict_proba(X_test)[:,1]
fpr, tpr, thresholds=roc_curve(y_test, y_pred_proba)
```

```
plt.plot([0,1], [0,1], '')
plt.plot(fpr, tpr, label='')
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.title('Logistic Regression ROC curve')
plt.show()
```

⮕

```python
from sklearn.metrics import roc_auc_score #area under the ROC curve
roc_auc_score(y_test, y_pred_proba)
```