

Sentiment analysis is a field dedicated to extracting subjective emotions and feelings from text. One common way to express negative or positive feelings. Written reviews are great datasets for doing sentiment analysis because they are easy to collect and can be used to train an algorithm.

You'll work with the IMDB dataset: a set of 50,000 highly polarized reviews from the Internet Movie Database. 25,000 reviews for training, each set consisting of 50% negative and 50% positive reviews.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from keras.datasets import imdb
```



The argument `num_words=7500` means you'll only keep the top 7500 most frequently occurring words in the dataset. This allows you to work with vector data of manageable size. The variables `train_data` and `test_data` are lists of reviews (each review is a sequence of words)....list of lists! `train_labels` and `test_labels` are lists of 0s and 1s, where 0 stands for negative and 1 for positive. We will try to convert list of lists into a dataframe. `train_data` and `test_data` contain a sequence (list) of words.

```
vocabulary=7500
```

```
# save np.load
#np_load_old = np.load

# modify the default parameters of np.load
#np.load = lambda *a,**k: np_load_old(*a, allow_pickle=True, **k)

# call load_data with allow_pickle implicitly set to true
#(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=vocabulary)

# restore np.load for future normal usage
#np.load = np_load_old
```

```
np.load.__defaults__=(None, True, True, 'ASCII')
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=vocabulary)
np.load.__defaults__=(None, False, True, 'ASCII')
```

Loads the data as a list of integers

```
print(train_data[0])
```



It shows that Tokenizing has been done allready

It shows that tokenizing has been done already

```
print(len(train_data[100]))
```



Reviews have already been converted in a sequence of indexed words. One-hot encode the lists to turn them into an array instance, turning the sequence [3, 5] into a 10,000-dimensional vector that would be all 0s except for indices 3 and 5.

We define a function called `vectorize_sequences` that takes 2 arguments: review and dimension of vocabulary.

```
def vectorize_sequences(sequences, dimension=10000):  
    results=np.zeros((len(sequences), dimension)) #matrix of 10000 columns and with rows  
    for i, sequence in enumerate(sequences):  
        results[i, sequence]=1  
    return results
```

```
X_train=vectorize_sequences(train_data)  
X_test=vectorize_sequences(test_data)
```

```
print(X_train[0])  
print(X_test[0])
```

Using TensorFlow backend.

```
X_train.shape
```

[1, 14, 22, 16, 43, 530, 973, 1622, 1385, 65, 458, 4468, 66, 3941, 4, 173, 3]

```
X_test.shape
```

158

We currently have a list of lists and will convert to a panda dataframe using the pandas DataFrame constructor.

```
X_train_df=pd.DataFrame(X_train)  
y_train_df=pd.DataFrame(train_labels, columns=['IMDB training labels'])  
  
X_test_df=pd.DataFrame(X_test)  
y_test_df=pd.DataFrame(test_labels, columns=['IMDB Testing labels'])
```

```
X_train_df.head()
```

[0. 1. 1. ... 0. 0. 0.]

```
[0. 1. 1. ... 0. 0. 0.]
```

```
y_train_df['IMDB training labels'].value_counts()
```

```
↳ (25000, 10000)
```

Dataset is balanced

we want to merge the training data and labels to make up 1 dataframe

```
train_df = pd.concat([X_train_df, y_train_df], axis=1) #combined training data and labels
```

```
test_df = pd.concat([X_test_df, y_test_df], axis=1) #combined training data and labels
```

```
train_df.head()
```

```
↳ (25000, 10000)
```

We will use a logistic regression classifier - categorizes data in 2 classes via Sigmoid

```
from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression()
classifier.fit(X_train_df, y_train_df)
score = classifier.score(X_test_df, y_test_df)
print("Accuracy:", score)
```



	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17
0	0.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0
1	0.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0	1.0	0.0	1.0	1.0	1.0	1.0	0.0
2	0.0	1.0	1.0	0.0	1.0	0.0	1.0	1.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0
3	0.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
4	0.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0

Logistics Regression gives us an accuracy of 86%

```
from sklearn.naive_bayes import MultinomialNB
```

```
model=MultinomialNB()
```

```
model.fit(X_train_df, y_train_df);
score=model.score(X_test_df, y_test_df)
print("Accuracy:", score)
```



	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17
0	0.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0
1	0.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0	1.0	0.0	1.0	1.0	1.0	1.0	0.0
2	0.0	1.0	1.0	0.0	1.0	0.0	1.0	1.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0

Naive Bayes - Multinomial gives us an accuracy of 84%

```
from sklearn.naive_bayes import GaussianNB
```

```
model=GaussianNB()
```

```
model.fit(X_train_df, y_train_df);
score=model.score(X_test_df, y_test_df)
print("Accuracy:", score)
```

```

/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:432:
FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py:724: Data
y = column_or_1d(y, warn=True)
Accuracy: 0.85852

```

Naive Bayes - Gaussian gives us an accuracy of 69% - not a surprise as multinomial assumption is known to be wrong. Gaussian assumption

Now we will look at Random Forests and Adaboost

```
from sklearn.ensemble import RandomForestClassifier
```

```
clf=RandomForestClassifier(n_estimators=100) #n_estimators is the number of trees in the forest
```

```
clf.fit(X_train_df,y_train_df)
```

```

/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py:724: Data
y = column_or_1d(y, warn=True)
Accuracy: 0.8424

```

```
y_pred_df=clf.predict(X_test_df)
from sklearn import metrics
print("Accuracy:",metrics.accuracy_score(y_test_df, y_pred_df))
```

```
↳ /usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py:724: Data
    y = column or 1d(y, warn=True)
```

```
from sklearn.ensemble import AdaBoostClassifier
clf = AdaBoostClassifier(n_estimators=50,learning_rate=1,random_state=0)
model = clf.fit(X_train_df, y_train_df)
```

```
print("Accuracy:",metrics.accuracy_score(y_test_df, y_pred_df))
```